

BA AI WORKFLOW TOOLKIT

A practical user-story grooming skill + reusable AI risk checklist

Designed for Business Analysts who want AI to strengthen their workflows without forfeiting human judgment and decision-making.

What this toolkit is

This toolkit packages a repeatable AI-assisted workflow for preparing user stories for grooming. It is intentionally more than a prompt: it defines inputs, steps, boundaries, review points, and an output structure.

The core idea is that you can use AI to help analyze complex and lengthy source material and expose gaps so the team can spend more meeting time on strategic decisions and less time in discovery.

What this toolkit is not

- A substitute for Business Analyst judgment.
- A mechanism for inventing missing business rules.
- A sizing or sprint-planning tool.
- Permission to place confidential or restricted company information into an unapproved AI environment.

Quick start

- 1. Adapt it:** Replace any terminology, meeting cadence, ticketing system, or output format with the way your team actually works. Or keep it as is. This was built on a proven structure so you may find that fewer adjustments are better.
- 2. Confirm your environment:** Make sure your organization permits allows use of the AI tool and data you intend to use.
- 3. Start with one story before processing batches:** Provide content from just one ticket along with visual/design, meeting notes or other reference materials. Process one story at a time until you are satisfied with your workflow, then start using it to process many user stories at a time.
- 4. Review the draft:** Validate criteria against source material. Resolve assumptions, mismatches, and dependencies with the right people.
- 5. Feed learning forward:** After grooming, carry unresolved questions into the next preparation cycle.

Skill: Weekly Story Grooming Prep

Public/adaptable version

Purpose

This skill turns raw work-item tickets and visual comps into meeting-ready user stories. The meeting should be able to focus on decision-making, technical risks/dependencies, and next steps rather than spending time discovering basic gaps.

Story sizing and sprint placement remain with the development team and are outside this skill.

Treat every output as a draft for human review. The value of the skill is to ultimately process a high volume of user stories with first-pass story writing and gap detection. It does not replace BA judgment about what is true, risky, or ready to commit to.

Before starting: data sensitivity pre-check

Before processing any ticket or comp, scan it for content that should not be shared with the AI tool or environment being used, including client-identifying information, financial data, real customer/employee personal data, vendor contract terms, credentials, or unreleased pricing, roadmap, or strategic initiatives.

If questionable content is present, stop rather than guessing whether it is safe to process.

Workflow

Process one story at a time so assumptions and gaps do not blur across stories.

Step 1: Intake

Confirm that both the ticket content and the visual comp are available. The comp may be an image, description, or approved link. If either is missing, flag it rather than drafting from incomplete information.

Step 2: Write the User Story Statement

Write a concise "As a / I want / So that" statement.

- Write from the point of view of the relevant end user.
- Keep each clause concise.
- Base the statement on the actual problem and intended solution, not a generic restatement of the ticket title.

Step 3: Draft Acceptance Criteria

Read the ticket and comp together.

- Base each criterion only on what is stated or reasonably inferable from the approved inputs.
- Cover the primary flow, at least one relevant edge case, and any error/empty state visible in the comp.

- If the comp and ticket conflict, call out the mismatch rather than silently choosing one.
- Order criteria in a way that makes review easy for your team. If visual layout matters, mirror the comp's visual sequence.

Suggested format:

Scenario 1: [Scenario title]

Verify that [specific, testable behavior].

Verify that [another specific, testable behavior].

Scenario 2: [Scenario title]

Verify that [specific, testable behavior].

Always include an Assumptions & Open Questions section. Anything inferred, unclear, or requiring team confirmation belongs there.

Step 4: Gap Check

Check the story for:

- Clear problem statement
- Visual comp or design reference
- Acceptance criteria
- Dependencies
- Unresolved questions or mismatches

If another team appears to be a dependency, flag it for confirmation. Do not claim that a cross-team dependency is resolved based on the ticket alone.

Step 5: Consolidate

Combine the user story statement, acceptance criteria, flagged gaps, and assumptions/open questions into one per-story summary. Confirm that nothing was dropped or materially altered during consolidation.

Step 6: Feed Forward

After the meeting, carry unresolved gaps and questions into the next preparation cycle. Mark something resolved only when a human confirms that it was resolved.

Output Format

[Story Title / Work Item ID]

User Story

As a [user], I want [goal], so that [benefit].

Acceptance Criteria

Scenario 1: [Scenario title]

Verify that [behavior].

Verify that [behavior].

Flagged Gaps

- ...

Assumptions & Open Questions

- ...

Verify Before Using the Output

- Do the acceptance criteria match the approved source material rather than a plausible guess?
- Is anything too vague to test?
- Are cross-team dependencies confirmed with the owning team?
- Is anything present that should not be shared with the meeting audience?
- Has a human made the final call on readiness, commitments, and next steps?

Reusable AI Workflow Risk Checklist

Before starting

- ☐ Confirm the AI tool/environment is approved for this data. Check with your organization if unsure.
- ☐ Scan inputs for restricted content: client-identifying information, financial data, real personal data, unreleased pricing/roadmap, credentials, or other confidential material.

Choose the simplest safe building block

- ☐ Is the step organizing, formatting, or summarizing known information? Use straightforward summarization.
- ☐ Does it require inferring intent, weighing tradeoffs, or working through incomplete information? Use reasoning deliberately.
- ☐ Does it require a fact that could be outdated or wrong? Verify it with an approved source/search capability.
- ☐ Does it require broad synthesis across many external sources? Use deeper research only when the task truly needs it.
- ☐ Default to the lightest capability that does the job.

Before sharing or acting on output

- ☐ Does every claim or criterion trace back to an actual input rather than an inference presented as fact?
- ☐ Are assumptions and open questions explicitly flagged?
- ☐ Are dependencies on other people or teams confirmed with them directly?
- ☐ Is confidential, restricted, or personal information present that should not be shared with this audience?
- ☐ Would you be comfortable if the output were reviewed by your manager or the people it concerns, as-is?
- ☐ Is there a policy question that should be confirmed with your organization?

Human decision points: do not delegate

- ☐ Final confirmation of cross-team or vendor dependencies
- ☐ Final sprint, timing, or delivery commitments
- ☐ Decisions involving confidential, regulated, or personal data

☐ Final go/no-go on what enters the meeting or gets sent onward

How to make the skill yours

User perspective: Who should the story be written for? The front-end customer, an internal user, an administrator, or another persona?

Inputs: What does your team rely on besides the ticket? Designs, process maps, API specs, transcripts, emails, research, or linked documentation?

Definition of 'meeting ready': What must be true before your team can have a productive grooming conversation?

Acceptance-criteria style: Use the format your team can review and test most easily. Do not adopt a format simply because AI defaults to it.

Known failure modes: Write explicit instructions for the mistakes you most want the skill to avoid, such as inventing rules, overlooking visual states, or assuming dependencies are resolved.

Human decision points: Name the decisions that remain with people. Make these boundaries visible rather than implicit.

A useful design test

For every step, ask: What is the AI allowed to infer? What must it flag instead? What source should the output trace back to? And where does a human make the final call?

Sharing note

This toolkit is meant to be adapted. Before using it at work, remove organization-specific examples, follow your employer's AI and data-handling policies, and validate the output before it enters a meeting, ticket, or downstream process.

Build the workflow. Keep the judgment.